

Published: 2018-10-21

[Table of Contents](#)

Wildcard certificates from Let's Encrypt with cert-manager and ingress-nginx on Google Kubernetes Engine

Goal

The goal of this guide is to have a simple web service running on a Google Kubernetes Engine cluster with wildcard certificates from Let's Encrypt and using the `ingress-nginx` as the Ingress controller. The certificates will be managed by `cert-manager`.

Prerequisites

- Working GKE cluster
 - running Kubernetes 1.10+
- Domain that you own, using Google Cloud DNS nameservers
 - this guide will use `example.xyz` in place of a real domain
 - if you have not yet configured your domain registrar for this, refer to the “Configure your domain registrar to use Google Cloud DNS’ nameservers” section below.
- Docker and the `gcloud` tools installed on your machine

This guide assumes you already know how to work with Pods, Deployments, Services, and Secrets on Kubernetes.

Trying out the simple web service locally

The app we'll be using, called `wildcard-demo`, is available on [Github](#) and [Docker Hub](#). All it does is return some request and host data in the response, which would normally be useful for apps that work with multiple subdomains.

Pull the image from Docker Hub and run it with the following commands:

```
$ docker pull john2x/wildcard-demo
$ docker run -p 8080:8080 john2x/wildcard-demo
```

With the server running, accessing `http://localhost:8080` or `http://foo.localhost:8080` should return something like the following:

```
$ curl http://localhost:8080 -i
HTTP/1.1 200 OK
Server: gunicorn/19.9.0
Date: Wed, 26 Sep 2018 05:59:11 GMT
Connection: close
Content-Type: text/plain
Content-Length: 193

remote address: 127.0.0.1
x-forwarded-for: None
```

```
hostname: localhost:8080
pod ip: None
pod name: None
node name: None
```

[Table of Contents](#)

Most of the information are not available since we are not in a cluster yet.

To stop the server, first find the container id of the image and then stop it with the `docker stop` command.

```
$ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED             STATUS
6d17c2e30642       john2x/wildcard-demo  "/start.sh"        3 minutes ago       Up
...
$ docker stop 6d17c2e30642
```

The source code of the app is on Github if you want to review the code before deploying it to your cluster, or if you want to build the image yourself.

Deploying the app to the cluster

Now that we've seen what the demo app does, we can deploy it to the cluster and access it from there.

But first let's create a new namespace for the demo app so everything is neat and organized.

```
$ kubectl create namespace wildcard-demo
namespace "wildcard-demo" created
...
$ kubectl get namespace
NAME                STATUS    AGE
default             Active   20m
kube-public         Active   20m
kube-system         Active   20m
wildcard-demo       Active   1m
```

With the namespace ready, let's take a look at our Deployment.

Listing 1: 01-deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: wildcard-demo
  namespace: wildcard-demo
  labels:
    app: wildcard-demo
spec:
  replicas: 2
  selector:
    matchLabels:
      app: wildcard-demo
  template:
    metadata:
      labels:
        app: wildcard-demo
    spec:
      containers:
```

```
- name: wildcard-demo
  image: john2x/wildcard-demo:latest
  imagePullPolicy: Always
  ports:
    - containerPort: 8080
      name: gunicorn
  resources:
    requests:
      cpu: 100m
      memory: 32Mi

  env:
    - name: POD_IP
      valueFrom:
        fieldRef:
          fieldPath: status.podIP
    - name: POD_NAME
      valueFrom:
        fieldRef:
          fieldPath: metadata.name
    - name: NODE_NAME
      valueFrom:
        fieldRef:
          fieldPath: spec.nodeName
```

[Table of Contents](#)

It's a fairly simple Deployment with some pod and node metadata exposed as environment variables.

Save the YAML above to a file and apply it with `kubectl apply -f <file>`.

Next we'll create a Service so we can access the app.

Listing 2: 02-service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: wildcard-demo
  namespace: wildcard-demo
spec:
  type: ClusterIP
  selector:
    app: wildcard-demo
  ports:
    - port: 80
      targetPort: gunicorn
      name: http
```

With the Deployment and Service ready, we should be able to test the app using port-forwarding. Get a pod name and run the following command:

```
$ kubectl port-forward svc/wildcard-demo 8080:80 --namespace=wildcard-demo
```

Accessing `http://localhost:8080` or `http://foo.localhost:8080` again should show a similar result as when we ran the image locally, only this time we'll get values for the pod and node details.

```
$ curl http://localhost:8080 -i
HTTP/1.1 200 OK
Server: gunicorn/19.9.0
```

```
Date: Wed, 26 Sep 2018 05:59:11 GMT
Connection: close
Content-Type: text/plain
Content-Length: 193
```

```
remote address: 127.0.0.1
x-forwarded-for: None
hostname: localhost:8080
pod ip: 10.8.0.8
pod name: wildcard-demo-76dd957877-84jqb
node name: gke-wildcard-demo--default-pool-e14a69d3-br01
```

[Table of Contents](#)

Installing ingress-nginx on the cluster

Now that we've tested that the web service is running and accessible, let's set up the Ingress so we can access it externally without having to do port-forwarding.

We'll be using the [ingress-nginx](https://kubernetes.github.io/ingress-nginx/deploy) Ingress controller, and we defer to the official deployment guide at <https://kubernetes.github.io/ingress-nginx/deploy>. You will need to run the commands under "Mandatory command" and "GCE - GKE". For reference, we mirror the commands here but it is recommended that you visit the official page in case there have been changes to the instructions.

```
$ kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/master/deploy/namespace "ingress-nginx" configured
deployment.extensions "default-http-backend" created
service "default-http-backend" created
configmap "nginx-configuration" created
configmap "tcp-services" created
configmap "udp-services" created
serviceaccount "nginx-ingress-serviceaccount" created
clusterrole.rbac.authorization.k8s.io "nginx-ingress-clusterrole" created
role.rbac.authorization.k8s.io "nginx-ingress-role" created
rolebinding.rbac.authorization.k8s.io "nginx-ingress-role-nisa-binding" created
clusterrolebinding.rbac.authorization.k8s.io "nginx-ingress-clusterrole-nisa-binding" created
deployment.extensions "nginx-ingress-controller" created

$ kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/master/deploy/service "ingress-nginx" created
```

Missing ClusterRoleBinding for cluster-admin ClusterRole

If you run into forbidden errors when running the commands like the one below, that probably means your Google account does not have the `cluster-admin` ClusterRole bound to it.

```
Error from server (Forbidden): error when creating "https://raw.githubusercontent.com/k
```

To bind your account to the `cluster-admin` ClusterRole, run the following command:

```
$ kubectl create clusterrolebinding cluster-admin-john --clusterrole=cluster-admin --use
```

Once your account has been bound, re-run the `ingress-nginx` set up commands.

Testing ingress-nginx's default HTTP backend

To verify that the Ingress controller is properly installed, we can visit the default HTTP backend that is included in the set up by visiting its LoadBalancer Service's external IP (I have redacted the last part of it).

[Table of Contents](#)

```
$ kubectl get service ingress-nginx --namespace=ingress-nginx
NAME                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)
ingress-nginx       LoadBalancer       10.11.244.17     35.188.164.XYZ   80:32336/TCP,443:31053/TCP

$ curl http://35.188.164.XYZ -i
HTTP/1.1 404 Not Found
Server: nginx/1.15.3
Date: Wed, 26 Sep 2018 06:06:02 GMT
Content-Type: text/plain; charset=utf-8
Content-Length: 21
Connection: keep-alive

default backend - 404
```

What we have so far

At this point we have the following pieces working:

1. Our wildcard-demo app deployed to its own Namespace, with a Service that's accessible only via port-forwarding.
2. The `ingress-nginx` Ingress controller (v0.19.0 at the time of writing), also deployed to its own Namespace, with a LoadBalancer Service accessible with an external IP address.

Add a subdomain record for ingress-nginx's default HTTP backend

Before we set up a subdomain for the wildcard-demo app, we should configure one for the default HTTP backend first since it's much simpler (being a single subdomain and all).

Configure your domain registrar to use Google Cloud DNS' nameservers

If you haven't updated your domain registrar to use Google Cloud DNS' nameservers yet, this would be a good time to do so.

The exact steps will be different for each registrar, but the overall process should be similar. The following steps will demonstrate the process for a domain registered via Namecheap. If you are using a different registrar, please refer to their respective documentation.

Create a Google Cloud DNS zone

Open Google Cloud Console and navigate to Networking > Network Services > Cloud DNS. Then create a Google Cloud DNS zone for your domain, like so:

Network services
Load balancing
Cloud DNS
Cloud CDN

Zone details

EDIT

ADD RECORD SET

DELETE ZONE

example-xyz

DNS name: example.xyz.

example.xyz is registered via namecheap.com

Record sets

Add record set

Delete record sets

DNS name ^	Type	TTL (seconds)	Data
example.xyz.	NS	21600	ns-cloud-a1.googledomains.com. ns-cloud-a2.googledomains.com. ns-cloud-a3.googledomains.com. ns-cloud-a4.googledomains.com.

Table of Contents

Note the nameservers listed in the NS record.

```
ns-cloud-a1.googledomains.com.
ns-cloud-a2.googledomains.com.
ns-cloud-a3.googledomains.com.
ns-cloud-a4.googledomains.com.
```

Change registrar's nameservers

Go to your domain's management page and change the nameservers to use the ones listed in your Google Cloud DNS zone.

Dashboard
Expiring / Expired
Domain List
Product List
Apps
Profile

example.xyz

Domain

Products

Sharing & Transfer

Advanced DNS

STATUS & VALIDITY

?

ACTIVE

Aug 31, 2018 - Sep 1, 2019

AUTO-RENEW

ADD YEARS

WhoisGuard™

?

PROTECTION

Aug 31, 2018 - Aug 31, 2019

AUTO-RENEW

ADD YEARS

SHOW DETAILS

PremiumDNS powered by VERISIGN

?

Enable PremiumDNS protection in order to switch your domain to our PremiumDNS platform. With our PremiumDNS platform, you get 100% DNS uptime and DDoS protection at the DNS level.

BUY NOW

NAMESERVERS

?

Custom DNS

ns-cloud-a1.googledomains.com

ns-cloud-a2.googledomains.com

ns-cloud-a3.googledomains.com

ns-cloud-a4.googledomains.com

ADD NAMESERVER

Wait and verify that the nameserver change have propagated

Use `dig` to check if the nameserver change have propagated. Note that propagation could take some time in some cases. (I

[Table of Contents](#)

```
$ dig NS example.xyz

; <<>> DiG 9.10.8-P1 <<>> NS example.xyz
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 46281
;; flags: qr rd ra; QUERY: 1, ANSWER: 4, AUTHORITY: 0, ADDITIONAL: 9

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 4096
;; QUESTION SECTION:
;example.xyz.                IN      NS

;; ANSWER SECTION:
example.xyz.                 3600    IN      NS      ns-cloud-a3.googledomains.com.
example.xyz.                 3600    IN      NS      ns-cloud-a2.googledomains.com.
example.xyz.                 3600    IN      NS      ns-cloud-a4.googledomains.com.
example.xyz.                 3600    IN      NS      ns-cloud-a1.googledomains.com.
```

Change the LoadBalancer's External IP address to static

Next is to change the external IP that was generated for the `ingress-nginx` LoadBalancer to a static one so it does not change. Open Google Cloud Console again and go to Networking > VPC Network > External IP addresses. Look for the correct external IP address and change its type from Ephemeral to Static.

Add the A record

Now that we have a static IP address and the domain nameservers are configured and propagated, we can add the A record for the default HTTP backend.

Go to your Google Cloud DNS zone and add an A record for your subdomain of choice using the external IP address of the `ingress-nginx` LoadBalancer Service.

Here we are using `default-http-backend.example.xyz` for the subdomain.

[Table of Contents](#)

The screenshot shows the 'Zone details' page for a DNS zone named 'example-xyz'. The page includes navigation links for 'EDIT', 'ADD RECORD SET', and 'DELETE ZONE'. It shows the DNS name 'example-xyz' and notes that it is registered via 'namecheap.com'. Under the 'Record sets' section, there are buttons for 'Add record set' and 'Delete record sets'. A table lists the current record sets:

DNS name	Type	TTL (seconds)	Data
example-xyz.	NS	21600	ns-cloud-a1.googledomains.com. ns-cloud-a2.googledomains.com. ns-cloud-a3.googledomains.com. ns-cloud-a4.googledomains.com.
example-xyz.	SOA	21600	ns-cloud-a1.googledomains.com. cloud-dns-hostmaster.google.com. 1 21600 3600 259200 300
default-http-backend.example-xyz.	A	300	35.188.164.XYZ

At the bottom, it says 'Equivalent REST' with a link.

Wait for about 5 minutes for the changes to propagate, and verify with `dig`.

```
$ dig A default-http-backend.example-xyz

; <<>> DiG 9.10.8-P1 <<>> A default-http-backend.example-xyz
;; global options: +cmd
;; Got answer:
;; ->HEADER<- opcode: QUERY, status: NOERROR, id: 29003
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 4, ADDITIONAL: 9

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 4096
;; QUESTION SECTION:
;default-http-backend.example-xyz. IN      A

;; ANSWER SECTION:
default-http-backend.example-xyz. 300 IN A      35.188.164.XYZ

;; AUTHORITY SECTION:
example-xyz.                1948 IN      NS      ns-cloud-a2.googledomains.com.
example-xyz.                1948 IN      NS      ns-cloud-a4.googledomains.com.
example-xyz.                1948 IN      NS      ns-cloud-a1.googledomains.com.
example-xyz.                1948 IN      NS      ns-cloud-a3.googledomains.com.

;; ADDITIONAL SECTION:
ns-cloud-a1.googledomains.com. 282117 IN A      216.239.32.106
ns-cloud-a1.googledomains.com. 83626 IN AAAA   2001:4860:4802:32::6a
ns-cloud-a2.googledomains.com. 282117 IN A      216.239.34.106
ns-cloud-a2.googledomains.com. 83626 IN AAAA   2001:4860:4802:34::6a
ns-cloud-a3.googledomains.com. 289144 IN A      216.239.36.106
ns-cloud-a3.googledomains.com. 83626 IN AAAA   2001:4860:4802:36::6a
ns-cloud-a4.googledomains.com. 273375 IN A      216.239.38.106
ns-cloud-a4.googledomains.com. 83626 IN AAAA   2001:4860:4802:38::6a

;; Query time: 274 msec
;; SERVER: 208.91.112.53#53(208.91.112.53)
;; WHEN: Wed Sep 26 15:03:39 +08 2018
;; MSG SIZE rcvd: 377
```

And try accessing the default backend via its new domain.


```
$ curl http://default-http-backend.example.xyz -i
HTTP/1.1 404 Not Found
Server: nginx/1.15.3
Date: Wed, 26 Sep 2018 07:07:25 GMT
Content-Type: text/plain; charset=utf-8
Content-Length: 21
Connection: keep-alive

default backend - 404
```

[Table of Contents](#)

Add a wildcard subdomain record for the wildcard-demo app

Now that we are able to access the default HTTP backend via a domain address, we want to be able to do the same for our wildcard-demo app. We want `http://wildcard-demo.example.xyz` or `http://foo.wildcard-demo.example.xyz` to point to the wildcard-demo app.

Create an Ingress for the wildcard-demo app

Apply the following resource.

Listing 3: 03-ingress-wildcard-demo.yaml

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: wildcard-demo-ingress
  namespace: wildcard-demo
  annotations:
    kubernetes.io/ingress.class: "nginx"
spec:
  rules:
    - host: wildcard-demo.example.xyz
      http:
        paths:
          - path: /
            backend:
              serviceName: wildcard-demo
              servicePort: http
    - host: "*.wildcard-demo.example.xyz"
      http:
        paths:
          - path: /
            backend:
              serviceName: wildcard-demo
              servicePort: http
```

Then confirm its creation:

```
$ kubectl get ingress --namespace=wildcard-demo
NAME                                HOSTS
wildcard-demo-ingress              wildcard-demo.example.xyz,*.wildcard-demo.example.xyz
```

We are not yet able to access the new Ingress at this point, since the hosts do not have a record in our DNS zone yet and the IP address will be routed to the default HTTP backend.

Add the A records for wildcard-demo

We need to add A records in our Google Cloud DNS zone, one for each of the hosts we specified in Ingress. Use the same external IP address for both A records.

[Table of Contents](#)

Zone details [EDIT](#) [ADD RECORD SET](#) [DELETE ZONE](#)

example-xyz [Registrar Setup](#)
 DNS name: **example-xyz.**
 example-xyz is registered via namecheap.com

Record sets
[Add record set](#) [Delete record sets](#)

<input type="checkbox"/> DNS name ^	Type	TTL (seconds)	Data
example-xyz.	NS	21600	ns-cloud-a1.googledomains.com. ns-cloud-a2.googledomains.com. ns-cloud-a3.googledomains.com. ns-cloud-a4.googledomains.com.
example-xyz.	SOA	21600	ns-cloud-a1.googledomains.com. cloud-dns-hostmaster.google.com. 1 21600 3600 259200 300
<input type="checkbox"/> default-http-handler.example-xyz.	A	300	35.188.164.XYZ
<input type="checkbox"/> wildcard-demo.example-xyz.	A	300	35.188.164.XYZ
<input type="checkbox"/> *.wildcard-demo.example-xyz.	A	300	35.188.164.XYZ

Wait for another 5 minutes, and verify with dig again.

```
$ dig A wildcard-demo.example.xyz

; <<>> DiG 9.10.8-P1 <<>> A wildcard-demo.example.xyz
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 782
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 4, ADDITIONAL: 9

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 4096
;; QUESTION SECTION:
;wildcard-demo.example.xyz. IN      A

;; ANSWER SECTION:
wildcard-demo.example.xyz. 300 IN      A      35.188.164.XYZ

...truncated...

$ dig A foo.wildcard-demo.example.xyz

; <<>> DiG 9.10.8-P1 <<>> A foo.wildcard-demo.example.xyz
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 18034
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 4, ADDITIONAL: 9

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 4096
;; QUESTION SECTION:
;foo.wildcard-demo.example.xyz. IN   A

;; ANSWER SECTION:
foo.wildcard-demo.example.xyz. 300 IN A      35.188.164.XYZ
```

...truncated...

[Table of Contents](#)

And finally test the new domain by accessing it.

```
$ curl http://wildcard-demo.example.xyz -i
HTTP/1.1 200 OK
Server: nginx/1.15.3
Date: Wed, 26 Sep 2018 07:26:11 GMT
Content-Type: text/plain
Content-Length: 217
Connection: keep-alive

remote address: 10.8.0.10
x-forwarded-for: 128.106.194.22
hostname: wildcard-demo.example.xyz
pod ip: 10.8.1.8
pod name: wildcard-demo-76dd957877-8cx7w
node name: gke-wildcard-demo--default-pool-e14a69d3-1t2m

$ curl http://foo.wildcard-demo.example.xyz -i
HTTP/1.1 200 OK
Server: nginx/1.15.3
Date: Wed, 26 Sep 2018 07:27:56 GMT
Content-Type: text/plain
Content-Length: 221
Connection: keep-alive

remote address: 10.8.0.10
x-forwarded-for: 128.106.194.22
hostname: foo.wildcard-demo.example.xyz
pod ip: 10.8.0.8
pod name: wildcard-demo-76dd957877-84jqb
node name: gke-wildcard-demo--default-pool-e14a69d3-br01
```

Note the x-Forwarded-For request header being set by ingress-nginx.

What we have so far

At this point we have the following additional pieces working:

1. An Ingress for the wildcard demo app.
2. A domain configured to use Google Cloud DNS nameservers.
3. A static external IP address for ingress-nginx.
4. The following A records in our Google Cloud DNS Zone:

```
default-http-backend.example.xyz. A    35.188.164.XYZ
wildcard-demo.example.xyz.         A    35.188.164.XYZ
*.wildcard-demo.example.xyz.       A    35.188.164.XYZ
```

Now that we are able to access our Services via their domain addresses, we are ready to request SSL certificates for them.

Installing cert-manager on the cluster

[cert-manager](#) is one of the more recent iterations for managing certificates on Kubernetes.

The recommended way to install cert-manager is via Helm, but that is beyond the scope of this guide. Refer to cert-manager's and/or Helm's official guides to install. Note that cert-manager also has instructions for installing using static manifests. Whichever method you follow to install, make sure you choose the one w

[Table of Contents](#)

Verifying your cert-manager install

By default, cert-manager will be installed in its own Namespace called `cert-manager`.

```
$ kubectl get deployment --namespace=cert-manager
NAME           DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
cert-manager   1         1         1            1           1m
```

Acquiring a “staging” certificate for ingress-nginx’s default HTTP backend

Before we attempt acquiring valid certificates, it's a good idea to try acquiring staging certificates first. This way we can debug any issues we run into without worrying about Let's Encrypt's rate limits.

Create a Google Cloud service account for managing Cloud DNS resources

In order for `cert-manager` to validate a domain via DNS-01 challenge, it needs to be able to add a TXT record to your Cloud DNS zone. We can authorize `cert-manager` to do this by creating a Google Cloud service account and storing it in a Secret.

Go to your Google Cloud Console and navigate to IAM & Admin > Service accounts. From there, create a new service account, assigning it a recognizable name and grant it access to the “DNS Administrator” role. When prompted to download the key, do so in the JSON format and save it to a file.

Now we need to upload the key stored in the JSON file in a Secret in the Namespace where the Issuer resource will be created (`ingress-nginx` in this case).

```
$ kubectl create secret generic clouddns-service-account --from-file=service-account-key.json
secret "clouddns-service-account" created

$ kubectl describe secret clouddns-service-account --namespace=ingress-nginx
Name:         clouddns-service-account
Namespace:    ingress-nginx
Labels:       <none>
Annotations:  <none>

Type: Opaque

Data
====
service-account-key.json: 2336 bytes
```

Create Issuer in ingress-nginx Namespace

An Issuer is a resource that issues certificates in a single Namespace. There's also a ClusterIssuer that can issue certificates across different Namespaces, but for this guide we will only use Issuer.

The following manifest creates an Issuer in the `ingress-nginx` Namespace and uses Let's Encrypt's staging API to request the certificates via a DNS-01 challenge using Google Cloud DNS as the provider.

Listing 4: 04-issuer-ingress-nginx-staging.yaml**Table of Contents**

```
apiVersion: certmanager.k8s.io/v1alpha1
kind: Issuer
metadata:
  name: letsencrypt-staging
  namespace: ingress-nginx
spec:
  acme:
    server: https://acme-staging-v02.api.letsencrypt.org/directory
    email: john@example.xyz

    # Name of a secret used to store the ACME account private key
    privateKeySecretRef:
      name: letsencrypt-staging

    # ACME DNS-01 provider configurations
    dns01:

      # Here we define a list of DNS-01 providers that can solve DNS challenges
      providers:
        - name: clouddns
          clouddns:
            # A secretKeyRef to a google cloud json service account
            serviceAccountSecretRef:
              name: clouddns-service-account
              key: service-account-key.json
            # The Google Cloud project in which to update the DNS zone
            project: example-project
```

Create Certificate resource

Once we have an Issuer we can create a Certificate.

Listing 5: 05-certificate-ingress-nginx-staging.yaml

```
apiVersion: certmanager.k8s.io/v1alpha1
kind: Certificate
metadata:
  name: default-http-backend-example-xyz-staging
  namespace: ingress-nginx
spec:
  secretName: default-http-backend-example-xyz-staging-tls
  issuerRef:
    name: letsencrypt-staging
    kind: Issuer
  commonName: default-http-backend.example.xyz
  dnsNames:
    - default-http-backend.example.xyz
  acme:
    config:
      - dns01:
          provider: clouddns
        domains:
          - default-http-backend.example.xyz
```

After a couple of minutes, the staging Certificate should be issued and stored in the Secret specified in the `secretName` specified above.

```
$ kubectl describe certificate default-http-backend-example-xyz-staging --namespace=ingress-nginx
Name:                default-http-backend-example-xyz-staging
Namespace:           ingress-nginx
Labels:              <none>
Annotations:         <truncated>
API Version:         certmanager.k8s.io/v1alpha1
Kind:                Certificate
Metadata:
  <truncated>
Spec:
  Acme:
    Config:
      Dns 01:
        Provider:  clouddns
      Domains:
        default-http-backend.example.xyz
    Common Name:   default-http-backend.example.xyz
    Dns Names:
      default-http-backend.example.xyz
    Issuer Ref:
      Kind:        Issuer
      Name:        letsencrypt-staging
    Secret Name:   default-http-backend-example-xyz-staging-tls
Status:
  Acme:
    Order:
      URL:  https://acme-staging-v02.api.letsencrypt.org/acme/order/7012226/8989674
  Conditions:
    Last Transition Time:  2018-09-26T12:44:31Z
    Message:              Certificate issued successfully
    Reason:               CertIssued
    Status:               True
    Type:                 Ready
    Last Transition Time:  <nil>
    Message:              Order validated
    Reason:               OrderValidated
    Status:               False
    Type:                 ValidateFailed
Events:
  Type    Reason          Age    From          Message
  ----    -
  Normal  CreateOrder     8m    cert-manager  Created new ACME order, attempting validation
  Normal  DomainVerified  6m    cert-manager  Domain "default-http-backend.example.xyz" is
  Normal  IssueCert       6m    cert-manager  Issuing certificate...
  Normal  CertObtained    6m    cert-manager  Obtained certificate from ACME server
  Normal  CertIssued      6m    cert-manager  Certificate issued successfully
```

[Table of Contents](#)

Use the Certificate in the default HTTP backend Ingress

We now have a TLS certificate stored in the `default-http-backend-example-xyz-staging-tls` Secret. Create a new Ingress for the default HTTP backend and use the certificate.

Listing 6: 06-ingress-default-http-backend.yaml

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: default-http-backend-ingress
  namespace: ingress-nginx
  annotations:
    kubernetes.io/ingress.class: "nginx"
```

```
spec:
  tls:
    - secretName: default-http-backend-example-xyz-staging-tls
      hosts:
        - default-http-backend.example.xyz
  rules:
    - host: default-http-backend.example.xyz
      http:
        paths:
          - path: /
            backend:
              serviceName: default-http-backend
              servicePort: 80
```

[Table of Contents](#)

Now when we access the default backend via http, we get a 308 Permanent Redirect to the https version thanks to ingress-nginx.

```
$ curl http://default-http-backend.example.xyz -i
HTTP/1.1 308 Permanent Redirect
Server: nginx/1.15.3
Date: Fri, 28 Sep 2018 01:51:39 GMT
Content-Type: text/html
Content-Length: 187
Connection: keep-alive
Location: https://default-http-backend.example.xyz/

<html>
<head><title>308 Permanent Redirect</title></head>
<body bgcolor="white">
<center><h1>308 Permanent Redirect</h1></center>
<hr><center>nginx/1.15.3</center>
</body>
</html>
```

Accessing via https gives us the familiar 404 response. Note the `-k` flag to disable SSL validation, since the staging certificates are not valid.

```
$ curl https://default-http-backend.example.xyz -i -k
HTTP/1.1 404 Not Found
Server: nginx/1.15.3
Date: Fri, 28 Sep 2018 01:54:10 GMT
Content-Type: text/plain; charset=utf-8
Content-Length: 21
Connection: keep-alive
Strict-Transport-Security: max-age=15724800; includeSubDomains

default backend - 404
```

Acquiring a “staging” wildcard certificate for the wildcard-demo app

The process is similar for wildcard certificates. First we create an Issuer in the `wildcard-demo` Namespace, then create the Certificate, and finally use the certificate in the Ingress.

Create Issuer in wildcard-demo Namespace

First we need to create the Secret to store the Cloud DNS service account key in the `wildcard-demo` Namespace.

```
$ kubectl create secret generic clouddns-service-account --from-file=service-account-key.json
secret "clouddns-service-account" created

$ kubectl describe secret clouddns-service-account --namespace=ingress-nginx
Name:         clouddns-service-account
Namespace:    wildcard-demo
Labels:       <none>
Annotations:  <none>

Type: Opaque

Data
====
service-account-key.json: 2336 bytes
```

[Table of Contents](#)

Then create the Issuer.

Listing 7: 07-issuer-wildcard-demo-staging.yaml

```
apiVersion: certmanager.k8s.io/v1alpha1
kind: Issuer
metadata:
  name: letsencrypt-staging
  namespace: wildcard-demo
spec:
  acme:
    server: https://acme-staging-v02.api.letsencrypt.org/directory
    email: john@example.xyz

    # Name of a secret used to store the ACME account private key
    privateKeySecretRef:
      name: letsencrypt-staging

    # ACME DNS-01 provider configurations
    dns01:

      # Here we define a list of DNS-01 providers that can solve DNS challenges
      providers:
        - name: clouddns
          clouddns:
            # A secretKeyRef to a google cloud json service account
            serviceAccountSecretRef:
              name: clouddns-service-account
              key: service-account-key.json
            # The Google Cloud project in which to update the DNS zone
            project: example-project
```

Create Certificate resource

Listing 8: 08-certificate-wildcard-demo-staging.yaml

```
apiVersion: certmanager.k8s.io/v1alpha1
kind: Certificate
metadata:
  name: wildcard-demo-example-xyz-staging
  namespace: wildcard-demo
spec:
  secretName: wildcard-demo-example-xyz-staging-tls
  issuerRef:
```



```

name: letsencrypt-staging
kind: Issuer
commonName: wildcard-demo.example.xyz
dnsNames:
- wildcard-demo.example.xyz
- "*.wildcard-demo.example.xyz"
acme:
  config:
    - dns01:
        provider: clouddns
      domains:
        - wildcard-demo.example.xyz
        - "*.wildcard-demo.example.xyz"

```

[Table of Contents](#)

Validating wildcard certificates takes longer, so wait a couple of minutes and check its status.

```

$ kubectl describe certificate wildcard-demo-example-xyz-staging --namespace=wildcard-d
Name:          wildcard-demo-example-xyz-staging
Namespace:     wildcard-demo
API Version:   certmanager.k8s.io/v1alpha1
Kind:          Certificate
Metadata:
  <truncated>
Spec:
  Acme:
    Config:
      Dns 01:
        Provider: clouddns
      Domains:
        wildcard-demo.example.xyz
        *.wildcard-demo.example.xyz
    Common Name: wildcard-demo.example.xyz
    Dns Names:
      wildcard-demo.example.xyz
      *.wildcard-demo.example.xyz
    Issuer Ref:
      Kind:      Issuer
      Name:      letsencrypt-staging
    Secret Name: wildcard-demo-example-xyz-staging-tls
Status:
  Acme:
    Order:
      URL: https://acme-staging-v02.api.letsencrypt.org/acme/order/7022259/9120998
  Conditions:
    Last Transition Time: 2018-09-28T02:13:13Z
    Message:             Certificate issued successfully
    Reason:              CertIssued
    Status:              True
    Type:                Ready
    Last Transition Time: <nil>
    Message:             Order validated
    Reason:              OrderValidated
    Status:              False
    Type:                ValidateFailed
Events:
  Type      Reason          Age          From          Message
  ----      -
  Normal    CreateOrder     7m           cert-manager   Created new ACME order, attempt
  Normal    DomainVerified  1m (x2 over 4m) cert-manager   Domain "wildcard-demo.example.
  Normal    IssueCert      1m           cert-manager   Issuing certificate...
  Normal    CertObtained   1m           cert-manager   Obtained certificate from ACME
  Normal    CertIssued     1m           cert-manager   Certificate issued successfully

```

Use the Certificate in the wildcard-demo Ingress

Finally, update the wildcard-demo Ingress to use the new certificate.

[Table of Contents](#)

Listing 9: 09-ingress-wildcard-demo-staging.yaml

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: wildcard-demo-ingress
  namespace: wildcard-demo
  annotations:
    kubernetes.io/ingress.class: "nginx"
spec:
  tls:
  - secretName: wildcard-demo-example-xyz-staging-tls
    hosts:
    - wildcard-demo.example.xyz
    - "*.wildcard-demo.example.xyz"
  rules:
  - host: wildcard-demo.example.xyz
    http:
      paths:
      - path: /
        backend:
          serviceName: wildcard-demo
          servicePort: http
  - host: "*.wildcard-demo.example.xyz"
    http:
      paths:
      - path: /
        backend:
          serviceName: wildcard-demo
          servicePort: http
```

Now we can access the app via https .

```
$ curl https://wildcard-demo.example.xyz -i -k
HTTP/1.1 200 OK
Server: nginx/1.15.3
Date: Fri, 28 Sep 2018 02:20:08 GMT
Content-Type: text/plain
Content-Length: 217
Connection: keep-alive
Strict-Transport-Security: max-age=15724800; includeSubDomains

remote address: 10.8.0.10
x-forwarded-for: 222.164.248.27
hostname: wildcard-demo.example.xyz
pod ip: 10.8.1.8
pod name: wildcard-demo-76dd957877-8cx7w
node name: gke-wildcard-demo--default-pool-e14a69d3-1t2m

$ curl https://foo.wildcard-demo.example.xyz -i -k
HTTP/1.1 200 OK
Server: nginx/1.15.3
Date: Fri, 28 Sep 2018 02:20:08 GMT
Content-Type: text/plain
Content-Length: 217
Connection: keep-alive
```

```
Strict-Transport-Security: max-age=15724800; includeSubDomains
```

```
remote address: 10.8.0.10
x-forwarded-for: 222.164.248.27
hostname: foo.wildcard-demo.example.xyz
pod ip: 10.8.1.8
pod name: wildcard-demo-76dd957877-8cx7w
node name: gke-wildcard-demo--default-pool-e14a69d3-1t2m
```

[Table of Contents](#)

Switching to valid Let's Encrypt certificates

We're almost done. We now have a working Issuer that we know doesn't run into issues when requesting for certificates from Let's Encrypt, and so it is safe to switch to using Let's Encrypt's production API's.

The process will be similar to how we got the staging certificates. First we need to create an Issuer in the Namespace that needs it, but this time the Issuer will be configured to use Let's Encrypt's production API. Then we create new Certificates using the production Issuer, and finally update the Ingresses to use the production Certificates.

Create production Issuers

For brevity, we will combine the two Issuers into one manifest file.

Listing 10: 10-issuer-wildcard-demo-prod.yaml

```
apiVersion: certmanager.k8s.io/v1alpha1
kind: Issuer
metadata:
  name: letsencrypt
  namespace: ingress-nginx
spec:
  acme:
    server: https://acme-v02.api.letsencrypt.org/directory
    email: john@example.xyz

    # Name of a secret used to store the ACME account private key
    privateKeySecretRef:
      name: letsencrypt

    # ACME DNS-01 provider configurations
    dns01:

      # Here we define a list of DNS-01 providers that can solve DNS challenges
      providers:
        - name: clouddns
          clouddns:
            # A secretKeyRef to a google cloud json service account
            serviceAccountSecretRef:
              name: clouddns-service-account
              key: service-account-key.json
            # The Google Cloud project in which to update the DNS zone
            project: example-project
---
apiVersion: certmanager.k8s.io/v1alpha1
kind: Issuer
metadata:
  name: letsencrypt
  namespace: wildcard-demo
```

```
spec:
  acme:
    server: https://acme-v02.api.letsencrypt.org/directory
    email: john@example.xyz

    # Name of a secret used to store the ACME account private key
    privateKeySecretRef:
      name: letsencrypt

    # ACME DNS-01 provider configurations
    dns01:

      # Here we define a list of DNS-01 providers that can solve DNS challenges
      providers:
        - name: clouddns
          clouddns:
            # A secretKeyRef to a google cloud json service account
            serviceAccountSecretRef:
              name: clouddns-service-account
              key: service-account-key.json
            # The Google Cloud project in which to update the DNS zone
            project: example-project
```

[Table of Contents](#)

Create production Certificates

Listing 11: 11-certificates-wildcard-demo-prod.yaml

```
apiVersion: certmanager.k8s.io/v1alpha1
kind: Certificate
metadata:
  name: default-http-backend-example-xyz
  namespace: ingress-nginx
spec:
  secretName: default-http-backend-example-xyz-tls
  issuerRef:
    name: letsencrypt
    kind: Issuer
  commonName: default-http-backend.example.xyz
  dnsNames:
    - default-http-backend.example.xyz
  acme:
    config:
      - dns01:
          provider: clouddns
        domains:
          - default-http-backend.example.xyz
---
apiVersion: certmanager.k8s.io/v1alpha1
kind: Certificate
metadata:
  name: wildcard-demo-example-xyz
  namespace: wildcard-demo
spec:
  secretName: wildcard-demo-example-xyz-tls
  issuerRef:
    name: letsencrypt
    kind: Issuer
  commonName: wildcard-demo.example.xyz
  dnsNames:
    - wildcard-demo.example.xyz
    - "*.wildcard-demo.example.xyz"
  acme:
```

```
config:
  - dns01:
      provider: clouddns
    domains:
      - wildcard-demo.example.xyz
      - "*.wildcard-demo.example.xyz"
```

[Table of Contents](#)

Use the new Certificates in our Ingresses

Listing 12: 12-ingress-wildcard-demo-prod.yaml

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: default-http-backend-ingress
  namespace: ingress-nginx
  annotations:
    kubernetes.io/ingress.class: "nginx"
spec:
  tls:
    - secretName: default-http-backend-example-xyz-tls
      hosts:
        - default-http-backend.example.xyz
  rules:
    - host: default-http-backend.example.xyz
      http:
        paths:
          - path: /
            backend:
              serviceName: default-http-backend
              servicePort: 80
---
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: wildcard-demo-ingress
  namespace: wildcard-demo
  annotations:
    kubernetes.io/ingress.class: "nginx"
spec:
  tls:
    - secretName: wildcard-demo-example-xyz-tls
      hosts:
        - wildcard-demo.example.xyz
        - "*.wildcard-demo.example.xyz"
  rules:
    - host: wildcard-demo.example.xyz
      http:
        paths:
          - path: /
            backend:
              serviceName: wildcard-demo
              servicePort: http
    - host: "*.wildcard-demo.example.xyz"
      http:
        paths:
          - path: /
            backend:
              serviceName: wildcard-demo
              servicePort: http
```

After applying the manifests above, we should now be able to access our backends via `https`, and this time their certificates will be valid.

Table of Contents

```
$ curl https://default-http-backend.example.xyz -i
HTTP/1.1 404 Not Found
Server: nginx/1.15.3
Date: Fri, 28 Sep 2018 06:40:10 GMT
Content-Type: text/plain; charset=utf-8
Content-Length: 21
Connection: keep-alive
Strict-Transport-Security: max-age=15724800; includeSubDomains

default backend - 404

$ curl https://foo.wildcard-demo.example.xyz -i
HTTP/1.1 200 OK
Server: nginx/1.15.3
Date: Fri, 28 Sep 2018 06:40:56 GMT
Content-Type: text/plain
Content-Length: 221
Connection: keep-alive
Strict-Transport-Security: max-age=15724800; includeSubDomains

remote address: 10.8.0.10
x-forwarded-for: 128.106.194.70
hostname: foo.wildcard-demo.example.xyz
pod ip: 10.8.1.8
pod name: wildcard-demo-76dd957877-8cx7w
node name: gke-wildcard-demo--default-pool-e14a69d3-1t2m
```

What we have so far

1. cert-manager (v0.5.0 at the time of writing) deployed to its own Namespace
2. The following cert-manager resources in the ingress-nginx Namespace
 - Issuer
 - Certificate (which in turn creates a Secret where the certificate key itself is stored)
3. The following cert-manager resources in the wildcard-demo Namespace
 - Issuer
 - Certificate (which in turn creates a Secret where the certificate key itself is stored)
4. Ingresses with their `tls` configured

Cleanup

Before we finish, we should delete the staging Certificates we acquired earlier. If you don't plan on using the staging Issuer again in the Namespace, go ahead and delete that as well.

```
$ kubectl delete certificate default-http-backend-example-xyz-staging --namespace=ingress-nginx
certificate.certmanager.k8s.io "default-http-backend-example-xyz-staging" deleted

$ kubectl delete certificate wildcard-demo-example-xyz-staging --namespace=wildcard-demo
certificate.certmanager.k8s.io "wildcard-demo-example-xyz-staging" deleted

$ kubectl delete issuer letsencrypt-staging --namespace=ingress-nginx
issuer.certmanager.k8s.io "letsencrypt-staging" deleted

$ kubectl delete issuer letsencrypt-staging --namespace=wildcard-demo
issuer.certmanager.k8s.io "letsencrypt-staging" deleted
```

Next steps

Table of Contents

I hope this guide helps clarify some of the confusion around how to set up wildcard certificates with Ingress controller.

A good exercise to go on from here would be to automate the process that was covered. It would be great if getting a new `wildcard-demo` instance in a separate Namespace, with its own subdomain and certificates automatically.

If you would like to learn more about the details of the various pieces at work, here are some links:

- ingress-nginx

official documentation

<https://kubernetes.github.io/ingress-nginx/>

- cert-manager

official documentation

<https://cert-manager.readthedocs.io/en/latest/>

using with ingress-nginx

<https://cert-manager.readthedocs.io/en/latest/tutorials/acme/securing-nginx-ingress-with-letsencrypt.html>

issuing certificates with DNS validation

<https://cert-manager.readthedocs.io/en/latest/tutorials/acme/dns-validation.html>

Issuer reference

<https://cert-manager.readthedocs.io/en/latest/reference/issuers.html>

ClusterIssuer reference

<https://cert-manager.readthedocs.io/en/latest/reference/clusterissuers.html>

- Let's Encrypt

official documentation

<https://letsencrypt.org/docs/>

staging environment

<https://letsencrypt.org/docs/staging-environment/>

DNS-01 validation

<https://docs.certifytheweb.com/docs/dns-validation.html>

Issues

Here are some of the issues I ran into while doing my own tests for this guide. These issues are probably just caused by mistakes on my part, but I'm listing them here in case someone finds them helpful.

cert-manager did not work when I tried a DNS zone for a subdomain instead of the domain

In my first attempt at creating the Cloud DNS zone, I didn't want to mix the test domains which will be created from this guide with my other active domains, and so I created a separate DNS zone for my test subdomain.

Network services

Load balancing

Cloud DNS

Cloud CDN

Zone details

EDIT

ADD RECORD SET

DELETE ZONE

test-example-xyz

DNS name: test.example.xyz.

example.xyz is registered via namecheap.com

Record sets

Add record set

Delete record sets

DNS name ^	Type	TTL (seconds)	Data
test.example.xyz.	NS	21600	ns-cloud-a1.googledomains.com. ns-cloud-a2.googledomains.com. ns-cloud-a3.googledomains.com. ns-cloud-a4.googledomains.com.
test.example.xyz.	SOA	21600	ns-cloud-a1.googledomains.com. cloud-dns-hostmaster.google.com. 1 21600 3600 259200 300

Equivalent REST

Table of Contents

This didn't play well with Let's Encrypt's verification process and so `cert-manager`'s Issuers couldn't acquire certificates.

I ended up deleting the extra DNS zone and just used the main DNS zone for my domain.

This is a [known issue](#).

Might need to modify `cert-manager`'s Deployment args to use `dns01-self-check-nameservers`

Another issue I ran into was with the DNS nameserver change not having been fully propagated yet when I started acquiring certificates. What added to the confusion was that when I ran `dig` on my machines, the DNS was good. But when checking `cert-manager`'s logs, I could see that it had trouble resolving my domains. I suspect this was due to not waiting at least 48 hours for the DNS changes to fully propagate. I was getting impatient at the time and so I ended up modifying `cert-manager`'s Deployment manifest and adding the `--dns01-self-check-nameservers` parameter to its container command.

```
...
spec:
  containers:
  - args:
    - --cluster-resource-namespace=$(POD_NAMESPACE)
    - --leader-election-namespace=$(POD_NAMESPACE)
    - --dns01-self-check-nameservers="1.1.1.1;53,8.8.8.8:53" # add this
  ...
```

Keywords: `kubernetes`, `google-cloud`, `lets-encrypt`

Modified: 2018-10-29 14:03:55 +08

Copyright (c) 2018 John Louis Del Rosario

[Emacs](#) 26.1 ([Org](#) mode 9.1.9)